

ECE - GY 6123

Introduction to Machine Learning

Term Project | Fall 2020

Hand Gesture Classifier using CNN and Transfer Learning

Mounting Google Drive where entire data and codes are created, stored and edited

In []:

```
from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

Importing required libraries and APIs

In []:

```
#Importing
import numpy as np
import tensorflow as tf
import keras
from keras.preprocessing.image import ImageDataGenerator
from keras.layers import Dense, Dropout, Activation, Add, MaxPooling2D, Conv2D, Flatten, BatchN
ormalization, MaxPool2D
from keras.models import Sequential
import matplotlib.pyplot as plt
from keras.preprocessing import image
```

Loading training and validation data from the designated folders in the Google Drive using Keras functions

In []:


```
#Getting Data
train_data = tf.keras.preprocessing.image_dataset_from_directory(
    "/content/drive/MyDrive/ML NYU Fall 2020/Machine Learning Term Project/Final Folder/T
raining_Data_Resized", labels='inferred', label_mode="categorical", class_names=['fantasti
c', 'hi-five', 'strength', 'thumbs-down', 'thumbs-up', 'victory', 'yolo'],
    color_mode="grayscale", batch_size=32, image_size=(28, 28), shuffle=True, seed=123,
    validation_split=0.15, subset="training"
)
val_data = tf.keras.preprocessing.image_dataset_from_directory(
    "/content/drive/MyDrive/ML NYU Fall 2020/Machine Learning Term Project/Final Folder/T
raining_Data_Resized", labels='inferred', label_mode="categorical", class_names=['fantast
ic', 'hi-five', 'strength', 'thumbs-down', 'thumbs-up', 'victory', 'yolo'],
    color_mode="grayscale", batch_size=32, image_size=(28, 28), shuffle=True, seed=123,
    validation_split=0.15, subset="validation"
)
```

Found 1402 files belonging to 7 classes.
Using 1192 files for training.
Found 1402 files belonging to 7 classes.
Using 210 files for validation.

Visualizing 9 random images from train data after pre-processing

In []:

```
#Visualizing Data
class_names = ['fantastic', 'hi-five', 'strength', 'thumbs-down', 'thumbs-up', 'victory', 'yolo']
plt.figure(figsize=(10, 10))
for images, labels in train_data.take(1):
    for i in range(9):
        ax = plt.subplot(3, 3, i + 1)
        tbp = image.img_to_array(images[i])
        plt.imshow(tf.squeeze(tbp), cmap = 'gray')
        plt.title(class_names[np.argmax(labels[i], axis=0)])
        plt.axis("off")
```



Model from Scratch

Creating model from the scratch

In []:

```
#Creating model
model = Sequential()
model.add(Conv2D(32, kernel_size=(3,3), activation = 'relu', input_shape=(28,28 ,1) ))
model.add(MaxPooling2D(pool_size = (2, 2)))
model.add(Conv2D(16, kernel_size = (3, 3), activation = 'relu'))
model.add(MaxPooling2D(pool_size = (2, 2)))
model.add(Flatten())
model.add(Dense(128, activation = 'relu'))
model.add(Dropout(0.20))
model.add(Dense(64, activation = 'relu'))
model.add(Dense(7, activation = 'softmax'))
```

In []:

```
#Model summary
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 26, 26, 32)	320
max_pooling2d (MaxPooling2D)	(None, 13, 13, 32)	0
conv2d_1 (Conv2D)	(None, 11, 11, 16)	4624
max_pooling2d_1 (MaxPooling2D)	(None, 5, 5, 16)	0
flatten (Flatten)	(None, 400)	0
dense (Dense)	(None, 128)	51328
dropout (Dropout)	(None, 128)	0
dense_1 (Dense)	(None, 64)	8256
dense_2 (Dense)	(None, 7)	455

Total params: 64,983
Trainable params: 64,983
Non-trainable params: 0

Creating callbacks; one for tracking time and the other for tracking accuracy [1]

In []:

```
from timeit import default_timer as timer

class TimingCallback(keras.callbacks.Callback):
    def __init__(self, logs={}):
        self.logs=[]
    def on_epoch_begin(self, epoch, logs={}):
        self.starttime = timer()
    def on_epoch_end(self, epoch, logs={}):
        self.logs.append(timer()-self.starttime)

cb = TimingCallback()
```

In []:

```
class myCallback(tf.keras.callbacks.Callback):
    def on_epoch_end(self, epoch, logs={}):
        if(logs.get('accuracy')>0.998):
            print("\nReached 99.8% accuracy so cancelling training!")
            self.model.stop_training = True

acc_callbacks = myCallback()
```

Compiling and training

In []:

```
#Training model
keras.backend.clear_session()
model.compile(optimizer='Adam', metrics=['accuracy'], loss='categorical_crossentropy')
fit = model.fit(train_data, validation_data=val_data, epochs=25, callbacks=[cb, acc_callbacks])
```

Epoch 1/25

38/38 [=====] - 3s 76ms/step - loss: 6.9843 - accuracy: 0.3070 - val_loss: 1.3435 - val_accuracy: 0.5190

Epoch 2/25

38/38 [=====] - 3s 77ms/step - loss: 0.9344 - accuracy: 0.6628 - val_loss: 0.5499 - val_accuracy: 0.8048

```

val_loss: 0.0040 - val_accuracy: 0.8540
Epoch 3/25
38/38 [=====] - 3s 78ms/step - loss: 0.4246 - accuracy: 0.8540 -
val_loss: 0.1666 - val_accuracy: 0.9619
Epoch 4/25
38/38 [=====] - 3s 76ms/step - loss: 0.1992 - accuracy: 0.9295 -
val_loss: 0.0674 - val_accuracy: 0.9810
Epoch 5/25
38/38 [=====] - 3s 70ms/step - loss: 0.0723 - accuracy: 0.9790 -
val_loss: 0.0165 - val_accuracy: 1.0000
Epoch 6/25
38/38 [=====] - 3s 67ms/step - loss: 0.0515 - accuracy: 0.9824 -
val_loss: 0.0078 - val_accuracy: 1.0000
Epoch 7/25
38/38 [=====] - 3s 68ms/step - loss: 0.0293 - accuracy: 0.9908 -
val_loss: 0.0046 - val_accuracy: 1.0000
Epoch 8/25
38/38 [=====] - 3s 67ms/step - loss: 0.0201 - accuracy: 0.9950 -
val_loss: 0.0065 - val_accuracy: 1.0000
Epoch 9/25
38/38 [=====] - 3s 70ms/step - loss: 0.0163 - accuracy: 0.9950 -
val_loss: 0.0050 - val_accuracy: 1.0000
Epoch 10/25
38/38 [=====] - 3s 70ms/step - loss: 0.0220 - accuracy: 0.9941 -
val_loss: 0.0046 - val_accuracy: 1.0000
Epoch 11/25
38/38 [=====] - 3s 69ms/step - loss: 0.0170 - accuracy: 0.9958 -
val_loss: 9.1314e-04 - val_accuracy: 1.0000
Epoch 12/25
38/38 [=====] - 3s 71ms/step - loss: 0.0125 - accuracy: 0.9958 -
val_loss: 0.0010 - val_accuracy: 1.0000
Epoch 13/25
38/38 [=====] - 3s 68ms/step - loss: 0.0165 - accuracy: 0.9950 -
val_loss: 0.0031 - val_accuracy: 1.0000
Epoch 14/25
38/38 [=====] - 3s 70ms/step - loss: 0.0141 - accuracy: 0.9958 -
val_loss: 5.3237e-04 - val_accuracy: 1.0000
Epoch 15/25
36/38 [=====>..] - ETA: 0s - loss: 0.0053 - accuracy: 1.0000
Reached 99.8% accuracy so cancelling training!
38/38 [=====] - 3s 70ms/step - loss: 0.0053 - accuracy: 1.0000 -
val_loss: 9.0092e-04 - val_accuracy: 1.0000

```

Printing time and plotting accuracy graphs

In []:

```

print(cb.logs)
print(sum(cb.logs))

```

```

[3.8499311109999326, 3.4155097029997705, 3.5635947370001304, 3.4466483050000534, 3.152695
532999587, 3.0459625500002403, 3.072302319999835, 2.978807604999929, 3.1337172190001183,
3.144875528999819, 3.1084275769999294, 3.1623346960000163, 3.070247517000098, 3.157531038
000343, 3.1343532389996653]
48.43693867899947

```

In []:

```

#Data for plots
acc = fit.history['accuracy']
val_acc = fit.history['val_accuracy']
loss = fit.history['loss']
val_loss = fit.history['val_loss']
epochs = range(1, len(loss) + 1)

```



In []:

```

#accuracy plot
plt.plot(epochs, acc, color='green', label='Training Accuracy')
plt.plot(epochs, val_acc, color='blue', label='Validation Accuracy')
plt.title('Training and Validation Accuracy')

```

```
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend()
plt.figure()
#loss plot
plt.plot(epochs, loss, color='pink', label='Training Loss')
plt.plot(epochs, val_loss, color='red', label='Validation Loss')
plt.title('Training and Validation Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.figure()
plt.show()
```



<Figure size 432x288 with 0 Axes>

In []:

```
#Saving model
model.save('/content/drive/MyDrive/ML NYU Fall 2020/Machine Learning Term Project/Final Folder/SignLanguageModelScratch.h5')
```

Testing the model on an unseen image

Indented block

In []:

```
#Predicting images
def Predict_image(dir,model):
    img = image.load_img(dir,color_mode='grayscale',target_size=(28,28,1))
    ar = image.img_to_array(img)
    print(ar.shape)
    plt.imshow(tf.squeeze(ar))
```

```

ar = np.expand_dims(ar,axis=0)
pred = model.predict(ar)
pred = pred.flatten()
for i in range(0,len(class_names)):
    print("{} => {:.2f}%".format(class_names[i],pred[i]*100))

```

```


Predict_image('/content/drive/MyDrive/ML NYU Fall 2020/Machine Learning Term Project/Final Folder/Test_Set_Resized/images302_resized.jpg',model)

```

```

(28, 28, 1)
fantastic => 0.76%
hi-five => 3.57%
strength => 7.44%
thumbs-down => 0.59%
thumbs-up => 16.51%
victory => 70.94%
yolo => 0.19%

```



In []:

```

import random, os
def Predict_Images(dir,model):
    plt.figure(figsize=(12, 12))
    class_names = ['fantastic', 'hi-five', 'strength', 'thumbs-down', 'thumbs-up', 'victory', 'yolo']
    count = 0
    for root, dirs, files in os.walk(dir, topdown=False):
        print(int(np.sqrt(len(files)))*int(np.sqrt(len(files))))
        for name in files:
            print(os.path.join(root, name))
            if np.random.rand()>0.5:
                ax = plt.subplot(3, 3, count+1)
                img = image.load_img(os.path.join(root, name), color_mode='grayscale', target_size=(28,28,1))
                ar = image.img_to_array(img)
                plt.imshow(tf.squeeze(ar), cmap = 'gray')
                ar = np.expand_dims(ar,axis=0)
                pred = np.argmax(model.predict(ar),axis=1)
                plt.title(class_names[pred[0]])
                count = count + 1
            if count == 9:
                break

```

```

Predict_Images("/content/drive/MyDrive/ML NYU Fall 2020/Machine Learning Term Project/Final Folder/Test_Set_Resized",model)

```


25

```

/content/drive/MyDrive/ML NYU Fall 2020/Machine Learning Term Project/Final Folder/Test_Set_Resized/34369987-man-hand-indicating-peace-victory-sign-the-v-sign-is-a-hand-gesture-in-which-the-index-and-middle-fi_resized.jpg
/content/drive/MyDrive/ML NYU Fall 2020/Machine Learning Term Project/Final Folder/Test_Set_Resized/6ip5qj7pT (1)_resized.jpg
/content/drive/MyDrive/ML NYU Fall 2020/Machine Learning Term Project/Final Folder/Test_Set_Resized/WIN_20201206_20_58_34_Pro_resized.jpg
/content/drive/MyDrive/ML NYU Fall 2020/Machine Learning Term Project/Final Folder/Test_Set_Resized/WIN_20201206_21_05_45_Pro_resized.jpg
/content/drive/MyDrive/ML NYU Fall 2020/Machine Learning Term Project/Final Folder/Test_S

```

/content/drive/MyDrive/ML NYU Fall 2020/Machine Learning Term Project/Final Folder/Test_Set_Resized/WIN_20201206_21_05_34_Pro_resized.jpg
/content/drive/MyDrive/ML NYU Fall 2020/Machine Learning Term Project/Final Folder/Test_Set_Resized/WIN_20201206_20_58_23_Pro_resized.jpg
/content/drive/MyDrive/ML NYU Fall 2020/Machine Learning Term Project/Final Folder/Test_Set_Resized/WIN_20201206_21_05_42_Pro_resized.jpg
/content/drive/MyDrive/ML NYU Fall 2020/Machine Learning Term Project/Final Folder/Test_Set_Resized/WIN_20201206_20_58_15_Pro_resized.jpg
/content/drive/MyDrive/ML NYU Fall 2020/Machine Learning Term Project/Final Folder/Test_Set_Resized/WIN_20201205_21_25_34_Pro_resized.jpg
/content/drive/MyDrive/ML NYU Fall 2020/Machine Learning Term Project/Final Folder/Test_Set_Resized/WIN_20201206_20_51_56_Pro_resized.jpg
/content/drive/MyDrive/ML NYU Fall 2020/Machine Learning Term Project/Final Folder/Test_Set_Resized/WIN_20201206_21_00_49_Pro_resized.jpg
/content/drive/MyDrive/ML NYU Fall 2020/Machine Learning Term Project/Final Folder/Test_Set_Resized/WIN_20201206_20_58_24_Pro_resized.jpg
/content/drive/MyDrive/ML NYU Fall 2020/Machine Learning Term Project/Final Folder/Test_Set_Resized/WIN_20201205_21_31_27_Pro_resized.jpg
/content/drive/MyDrive/ML NYU Fall 2020/Machine Learning Term Project/Final Folder/Test_Set_Resized/WIN_20201206_21_09_13_Pro_resized.jpg
/content/drive/MyDrive/ML NYU Fall 2020/Machine Learning Term Project/Final Folder/Test_Set_Resized/images270_resized.jpg
/content/drive/MyDrive/ML NYU Fall 2020/Machine Learning Term Project/Final Folder/Test_Set_Resized/images428_resized.jpg
/content/drive/MyDrive/ML NYU Fall 2020/Machine Learning Term Project/Final Folder/Test_Set_Resized/images302_resized.jpg
/content/drive/MyDrive/ML NYU Fall 2020/Machine Learning Term Project/Final Folder/Test_Set_Resized/images122_resized.jpg
/content/drive/MyDrive/ML NYU Fall 2020/Machine Learning Term Project/Final Folder/Test_Set_Resized/images427_resized.jpg



Following three cells are used to plot the output of intermediate convolution layers [2]

In []:

```
# Directory with our training horse pictures
fantastic_dir = os.path.join('/content/drive/MyDrive/ML NYU Fall 2020/Machine Learning Term Project/Final Folder/Training_Data_Resized/fantastic')
hi5_dir = os.path.join('/content/drive/MyDrive/ML NYU Fall 2020/Machine Learning Term Project/Final Folder/Training_Data_Resized/hi-five')
strength_dir = os.path.join('/content/drive/MyDrive/ML NYU Fall 2020/Machine Learning Term Project/Final Folder/Training_Data_Resized/strength')
thumbs_down_dir = os.path.join('/content/drive/MyDrive/ML NYU Fall 2020/Machine Learning Term Project/Final Folder/Training_Data_Resized/thumbs-down')
thumbs_up_dir = os.path.join('/content/drive/MyDrive/ML NYU Fall 2020/Machine Learning Term Project/Final Folder/Training_Data_Resized/thumbs-up')
victory_dir = os.path.join('/content/drive/MyDrive/ML NYU Fall 2020/Machine Learning Term Project/Final Folder/Training_Data_Resized/victory')
yolo_dir = os.path.join('/content/drive/MyDrive/ML NYU Fall 2020/Machine Learning Term Project/Final Folder/Training_Data_Resized/yolo')

fantastic_names = os.listdir(fantastic_dir)
hi5_names = os.listdir(hi5_dir)
strength_names = os.listdir(strength_dir)
thumbs_down_names = os.listdir(thumbs_down_dir)
thumbs_up_names = os.listdir(thumbs_up_dir)
victory_names = os.listdir(victory_dir)
yolo_names = os.listdir(yolo_dir)
```

In []:

```
from tensorflow.keras.preprocessing.image import img_to_array, load_img
# Let's define a new Model that will take an image as input, and will output
# intermediate representations for all layers in the previous model after
# the first.
successive_outputs = [layer.output for layer in model.layers[1:]]
#visualization_model = Model(img_input, successive_outputs)
visualization_model = tf.keras.models.Model(inputs = model.input, outputs = successive_outputs)
# Let's prepare a random input image from the training set.
fantastic_img_files = [os.path.join(fantastic_dir, f) for f in fantastic_names]
hi5_img_files = [os.path.join(hi5_dir, f) for f in hi5_names]
strength_img_files = [os.path.join(strength_dir, f) for f in strength_names]
thumbs_down_img_files = [os.path.join(thumbs_down_dir, f) for f in thumbs_down_names]
thumbs_up_img_files = [os.path.join(thumbs_up_dir, f) for f in thumbs_up_names]
victory_img_files = [os.path.join(victory_dir, f) for f in victory_names]
yolo_img_files = [os.path.join(yolo_dir, f) for f in yolo_names]
img_path = random.choice(fantastic_img_files + hi5_img_files + strength_img_files + thumbs_down_img_files + thumbs_up_img_files + victory_img_files + yolo_img_files)

img = load_img(img_path, color_mode='grayscale', target_size=(28,28,1)) # this is a PIL image
x = img_to_array(img) # Numpy array with shape (150, 150, 3)

x = x.reshape((1,) + x.shape) # Numpy array with shape (1, 150, 150, 1)

# Rescale by 1/255
#x /= 255

# Let's run our image through our network, thus obtaining all
# intermediate representations for this image.
successive_feature_maps = visualization_model.predict(x)

# These are the names of the layers, so can have them as part of our plot
layer_names = [layer.name for layer in model.layers[1:]]
```

In []:

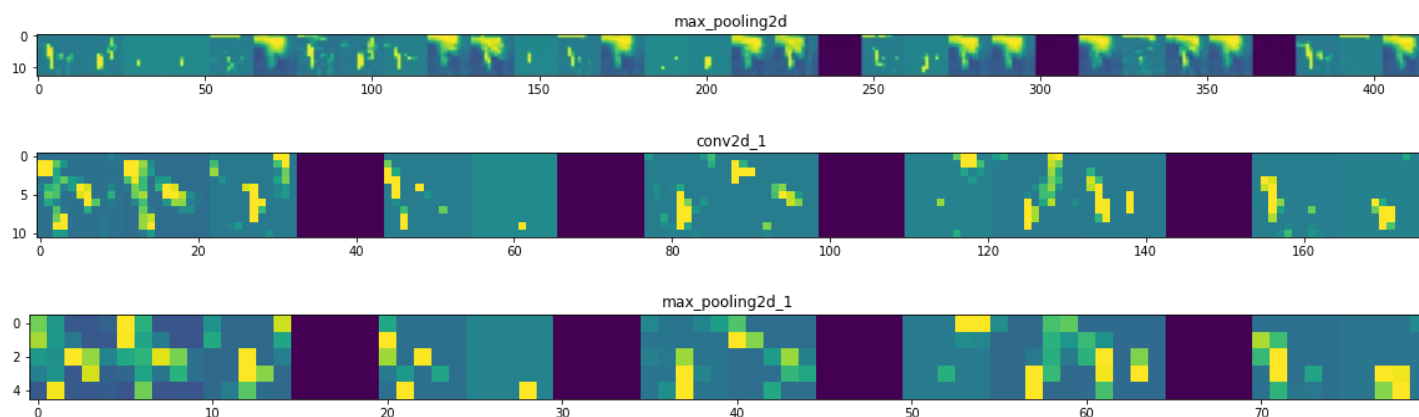
```
for layer_name, feature_map in zip(layer_names, successive_feature_maps):
```



```

if len(feature_map.shape) == 4:
    # Just do this for the conv / maxpool layers, not the fully-connected layers
    n_features = feature_map.shape[-1] # number of features in feature map
    # The feature map has shape (1, size, size, n_features)
    size = feature_map.shape[1]
    # We will tile our images in this matrix
    display_grid = np.zeros((size, size * n_features))
    for i in range(n_features):
        # Postprocess the feature to make it visually palatable
        x = feature_map[0, :, :, i]
        x -= x.mean()
        x /= x.std()
        x *= 64
        x += 128
        x = np.clip(x, 0, 255).astype('uint8')
        # We'll tile each filter into this big horizontal grid
        display_grid[:, i * size : (i + 1) * size] = x
    # Display the grid
    scale = 20. / n_features
    plt.figure(figsize=(scale * n_features, scale))
    plt.title(layer_name)
    plt.grid(False)
    plt.imshow(display_grid, aspect='auto', cmap='viridis')

```



Model using Transfer Learning

Loading fashion MNIST model saved in the directory [3]

In []:

```

from keras.models import load_model
#Loading MNIST model as base
model = load_model('/content/drive/MyDrive/ML NYU Fall 2020/Machine Learning Term Project /Final Folder/mnistModel')
model.layers[0].trainable = False
model.layers[1].trainable = False
model.summary()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 26, 26, 64)	640
max_pooling2d (MaxPooling2D)	(None, 13, 13, 64)	0
flatten (Flatten)	(None, 10816)	0
dense (Dense)	(None, 128)	1384576
dense_1 (Dense)	(None, 10)	1290

Total params: 1,386,506

Trainable params: 1,385,866

Non-trainable params: 640

Importing layers from fashion MNIST model and adding a few custom layers

In []:

```
#Creating new model
TL_Model = Sequential()
TL_Model.add(model.layers[0])
TL_Model.add(model.layers[1])
TL_Model.add(Conv2D(16, kernel_size = (3, 3), activation = 'relu',name='c2'))
TL_Model.add(MaxPooling2D(pool_size = (2, 2),name = 'm2'))
TL_Model.add(Flatten())
TL_Model.add(Dense(64,activation = 'relu'))
TL_Model.add(Dropout(0.20))
TL_Model.add(Dense(7,activation='softmax'))
```

In []:

```
TL_Model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 26, 26, 64)	640

max_pooling2d (MaxPooling2D)	(None, 13, 13, 64)	0

c2 (Conv2D)	(None, 11, 11, 16)	9232

m2 (MaxPooling2D)	(None, 5, 5, 16)	0

flatten (Flatten)	(None, 400)	0

dense (Dense)	(None, 64)	25664

dropout (Dropout)	(None, 64)	0

dense_1 (Dense)	(None, 7)	455
=====		
Total params: 35,991		
Trainable params: 35,351		
Non-trainable params: 640		

In []:

```
class TimingCallback(keras.callbacks.Callback):
    def __init__(self, logs={}):
        self.logs=[]
    def on_epoch_begin(self, epoch, logs={}):
        self.starttime = timer()
    def on_epoch_end(self, epoch, logs={}):
        self.logs.append(timer()-self.starttime)

cb1 = TimingCallback()
```

Compiling and training

In []:

```
#Training the model
keras.backend.clear_session()
TL_Model.compile(optimizer='Adam', metrics=['accuracy'], loss='categorical_crossentropy'
)
fit = TL_Model.fit(train_data,validation_data=val_data,epochs=25, callbacks=[cb1, acc_cal
lbacks])
```

```

38/38 [=====] - 2s 67ms/step - loss: 2.3392 - accuracy: 0.5443 -
val_loss: 0.4028 - val_accuracy: 0.8762
Epoch 2/25
38/38 [=====] - 2s 63ms/step - loss: 0.2778 - accuracy: 0.9111 -
val_loss: 0.0479 - val_accuracy: 0.9952
Epoch 3/25
38/38 [=====] - 2s 63ms/step - loss: 0.0965 - accuracy: 0.9723 -
val_loss: 0.0220 - val_accuracy: 0.9905
Epoch 4/25
38/38 [=====] - 2s 64ms/step - loss: 0.0600 - accuracy: 0.9799 -
val_loss: 0.0054 - val_accuracy: 1.0000
Epoch 5/25
38/38 [=====] - 2s 62ms/step - loss: 0.0381 - accuracy: 0.9866 -
val_loss: 0.0023 - val_accuracy: 1.0000
Epoch 6/25
36/38 [=====>..] - ETA: 0s - loss: 0.0145 - accuracy: 0.9983
Reached 99.8% accuracy so cancelling training!
38/38 [=====] - 2s 62ms/step - loss: 0.0145 - accuracy: 0.9983 -
val_loss: 0.0016 - val_accuracy: 1.0000

```

Printing time and plotting accuracy curves

In []:

```

print(cb1.logs)
print(sum(cb1.logs))

```

```

[3.337741518999792, 2.804591089000496, 2.8189245540006596, 2.861063084000307, 2.804527588
000383, 2.7703599050000776]
17.397207739001715

```

In []:

```

#Data for plots
acc = fit.history['accuracy']
val_acc = fit.history['val_accuracy']
loss = fit.history['loss']
val_loss = fit.history['val_loss']
epochs = range(1, len(loss) + 1)

```


In []:



```

#accuracy plot
plt.plot(epochs, acc, color='green', label='Training Accuracy')
plt.plot(epochs, val_acc, color='blue', label='Validation Accuracy')
plt.title('Training and Validation Accuracy')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend()
plt.figure()

#loss plot
plt.plot(epochs, loss, color='pink', label='Training Loss')
plt.plot(epochs, val_loss, color='red', label='Validation Loss')
plt.title('Training and Validation Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.figure()
plt.show()

```





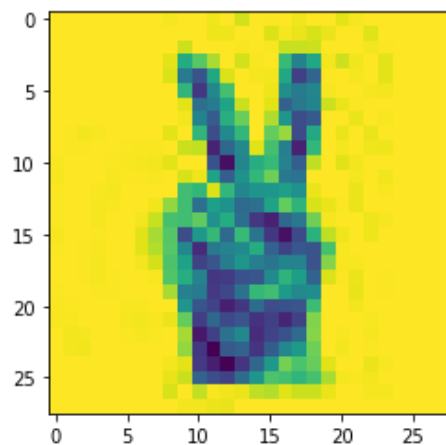
<Figure size 432x288 with 0 Axes>

Performance of transfer learning-based model on unseen images

In []:

```
Predict_image('/content/drive/MyDrive/ML NYU Fall 2020/Machine Learning Term Project/Final Folder/Test_Set_Resized/images302_resized.jpg', TL_Model)
```

```
(28, 28, 1)
fantastic => 0.00%
hi-five => 0.00%
strength => 0.15%
thumbs-down => 0.00%
thumbs-up => 0.00%
victory => 89.23%
yolo => 10.62%
```





In []:

```
Predict_Images("/content/drive/MyDrive/ML NYU Fall 2020/Machine Learning Term Project/Final Folder/Test_Set_Resized", TL_Model)
```

```
25
/content/drive/MyDrive/ML NYU Fall 2020/Machine Learning Term Project/Final Folder/Test_Set_Resized/34369987-man-hand-indicating-peace-victory-sign-the-v-sign-is-a-hand-gesture-in-which-the-index-and-middle-fi_resized.jpg
/content/drive/MyDrive/ML NYU Fall 2020/Machine Learning Term Project/Final Folder/Test_Set_Resized/6ip5qj7pT (1)_resized.jpg
/content/drive/MyDrive/ML NYU Fall 2020/Machine Learning Term Project/Final Folder/Test_Set_Resized/WIN_20201206_20_58_34_Pro_resized.jpg
/content/drive/MyDrive/ML NYU Fall 2020/Machine Learning Term Project/Final Folder/Test_Set_Resized/WIN_20201206_21_05_45_Pro_resized.jpg
```

et_Resized/WIN_20201206_21_05_34_Pro_resized.jpg
/content/drive/MyDrive/ML NYU Fall 2020/Machine Learning Term Project/Final Folder/Test_S
et_Resized/WIN_20201206_20_58_23_Pro_resized.jpg
/content/drive/MyDrive/ML NYU Fall 2020/Machine Learning Term Project/Final Folder/Test_S
et_Resized/WIN_20201206_21_05_42_Pro_resized.jpg
/content/drive/MyDrive/ML NYU Fall 2020/Machine Learning Term Project/Final Folder/Test_S
et_Resized/WIN_20201206_20_58_15_Pro_resized.jpg
/content/drive/MyDrive/ML NYU Fall 2020/Machine Learning Term Project/Final Folder/Test_S
et_Resized/WIN_20201205_21_25_34_Pro_resized.jpg
/content/drive/MyDrive/ML NYU Fall 2020/Machine Learning Term Project/Final Folder/Test_S
et_Resized/WIN_20201206_20_51_56_Pro_resized.jpg
/content/drive/MyDrive/ML NYU Fall 2020/Machine Learning Term Project/Final Folder/Test_S
et_Resized/WIN_20201206_21_00_49_Pro_resized.jpg
/content/drive/MyDrive/ML NYU Fall 2020/Machine Learning Term Project/Final Folder/Test_S
et_Resized/WIN_20201206_20_58_24_Pro_resized.jpg
/content/drive/MyDrive/ML NYU Fall 2020/Machine Learning Term Project/Final Folder/Test_S
et_Resized/WIN_20201205_21_31_27_Pro_resized.jpg
/content/drive/MyDrive/ML NYU Fall 2020/Machine Learning Term Project/Final Folder/Test_S
et_Resized/WIN_20201206_21_09_13_Pro_resized.jpg
/content/drive/MyDrive/ML NYU Fall 2020/Machine Learning Term Project/Final Folder/Test_S
et_Resized/images270_resized.jpg
/content/drive/MyDrive/ML NYU Fall 2020/Machine Learning Term Project/Final Folder/Test_S
et_Resized/images428_resized.jpg
/content/drive/MyDrive/ML NYU Fall 2020/Machine Learning Term Project/Final Folder/Test_S
et_Resized/images302_resized.jpg
/content/drive/MyDrive/ML NYU Fall 2020/Machine Learning Term Project/Final Folder/Test_S
et_Resized/images122_resized.jpg
/content/drive/MyDrive/ML NYU Fall 2020/Machine Learning Term Project/Final Folder/Test_S
et_Resized/images427_resized.jpg
/content/drive/MyDrive/ML NYU Fall 2020/Machine Learning Term Project/Final Folder/Test_S
et_Resized/images592_resized.jpg
/content/drive/MyDrive/ML NYU Fall 2020/Machine Learning Term Project/Final Folder/Test_S
et_Resized/images389_resized.jpg





Displaying output of convolutional layers for transfer learning-based models

In []:

```
successive_outputs = [layer.output for layer in TL_Model.layers[2:]]
#visualization_model = Model(img_input, successive_outputs)
visualization_model = tf.keras.models.Model(inputs = TL_Model.input, outputs = successive_outputs)

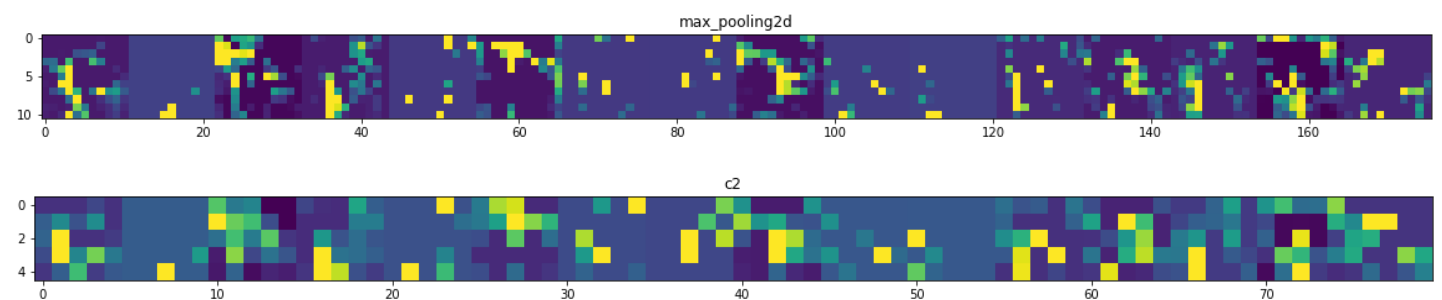
img = load_img(img_path, color_mode='grayscale', target_size=(28,28,1)) # this is a PIL image
x = img_to_array(img) # Numpy array with shape (150, 150, 3)

x = x.reshape((1,) + x.shape) # Numpy array with shape (1, 150, 150, 1)
successive_feature_maps = visualization_model.predict(x)

# These are the names of the layers, so can have them as part of our plot
layer_names = [layer.name for layer in TL_Model.layers[1:]]
```

In []:

```
for layer_name, feature_map in zip(layer_names, successive_feature_maps):
    if len(feature_map.shape) == 4:
        # Just do this for the conv / maxpool layers, not the fully-connected layers
        n_features = feature_map.shape[-1] # number of features in feature map
        # The feature map has shape (1, size, size, n_features)
        size = feature_map.shape[1]
        # We will tile our images in this matrix
        display_grid = np.zeros((size, size * n_features))
        for i in range(n_features):
            # Postprocess the feature to make it visually palatable
            x = feature_map[0, :, :, i]
            x -= x.mean()
            x /= x.std()
            x *= 64
            x += 128
            x = np.clip(x, 0, 255).astype('uint8')
            # We'll tile each filter into this big horizontal grid
            display_grid[:, i * size : (i + 1) * size] = x
        # Display the grid
        scale = 20. / n_features
        plt.figure(figsize=(scale * n_features, scale))
        plt.title(layer_name)
        plt.grid(False)
        plt.imshow(display_grid, aspect='auto', cmap='viridis')
```



References

- 1 'Introduction to TensorFlow for Artificial Intelligence, Machine Learning, and Deep Learning - Home | Coursera', <https://www.coursera.org/learn/introduction-tensorflow/ungradedLab/twxV5/lab>, accessed December 2020
- 2 'Course 1 - Part 8 - Lesson 2 - Notebook.ipynb - Colaboratory',

[https://colab.research.google.com/github/Imoroney/dlaicourse/blob/master/Course 1 - Part 8 - Lesson 2 - Notebook.ipynb#scrollTo=-5tES8rXFjux](https://colab.research.google.com/github/Imoroney/dlaicourse/blob/master/Course%201%20-%20Part%208%20-%20Lesson%202%20-%20Notebook.ipynb#scrollTo=-5tES8rXFjux), accessed December 2020

- 3 'Fashion MNIST | Kaggle', <https://www.kaggle.com/zalando-research/fashionmnist>, accessed December 2020